

Practical Exam: Grocery Store Sales

FoodYum is a grocery store chain that is based in the United States.

Food Yum sells items such as produce, meat, dairy, baked goods, snacks, and other household food staples.

As food costs rise, FoodYum wants to make sure it keeps stocking products in all categories that cover a range of prices to ensure they have stock for a broad range of customers.

Data

The data is available in the table `products`.

The dataset contains records of customers for their last full year of the loyalty program.

Column Name	Criteria
product_id	Nominal. The unique identifier of the product. Missing values are not possible due to the database structure.
product_type	Nominal. The product category type of the product, one of 5 values (Produce, Meat, Dairy, Bakery, Snacks). Missing values should be replaced with "Unknown".
brand	Nominal. The brand of the product. One of 7 possible values. Missing values should be replaced with "Unknown".
weight	Continuous. The weight of the product in grams. This can be any positive value, rounded to 2 decimal places. Missing values should be replaced with the overall median weight.
price	Continuous. The price the product is sold at, in US dollars. This can be any positive value, rounded to 2 decimal places. Missing values should be replaced with the overall median price.
average_units_sold	Discrete. The average number of units sold each month. This can be any positive integer value. Missing values should be replaced with 0.
year_added	Nominal. The year the product was first added to FoodYum stock. Missing values should be replaced with 2022.
stock_location	Nominal. The location that stock originates. This can be one of four warehouse locations, A, B, C or D Missing values should be replaced with "Unknown".

Task 1

In 2022 there was a bug in the product system. For some products that were added in that year, the `year_added` value was not set in the data. As the year the product was added may have an impact on the price of the product, this is important information to have.

Write a query to determine how many products have the `year_added` value missing. Your output should be a single column, `missing_year`, with a single row giving the number of missing values.

 Unknown integration DataFrame as `missing_year`

```
-- Write your query for task 1 in this cell
SELECT
    COUNT(*) AS missing_year
FROM products
WHERE year_added IS NULL;
```

index	...	↑↓	missing_year	...
			0	
Rows: 1				

 Expand

Task 2

Given what you know about the year added data, you need to make sure all of the data is clean before you start your analysis. The table below shows what the data should look like.

Write a query to ensure the product data matches the description provided. Do not update the original table.

Column Name	Criteria
product_id	Nominal. The unique identifier of the product. Missing values are not possible due to the database structure.
product_type	Nominal. The product category type of the product, one of 5 values (Produce, Meat, Dairy, Bakery, Snacks). Missing values should be replaced with "Unknown".
brand	Nominal. The brand of the product. One of 7 possible values. Missing values should be replaced with "Unknown".
weight	Continuous. The weight of the product in grams. This can be any positive value, rounded to 2 decimal places. Missing values should be replaced with the overall median weight.
price	Continuous. The price the product is sold at, in US dollars. This can be any positive value, rounded to 2 decimal places. Missing values should be replaced with the overall median price.
average_units_sold	Discrete. The average number of units sold each month. This can be any positive integer value. Missing values should be replaced with 0.
year_added	Nominal. The year the product was first added to FoodYum stock. Missing values should be replaced with last year (2022).
stock_location	Nominal. The location that stock originates. This can be one of four warehouse locations, A, B, C or D Missing values should be replaced with "Unknown".

Unknown integration DataFrame as c

`-- Write your query for task 2 in this cell`

```

WITH ranked_weights AS (
  SELECT
    weight,
    ROW_NUMBER() OVER (ORDER BY
      CASE
        WHEN weight ~ E'^\\d+\\.?\\d*\\s*grams$' THEN CAST(REPLACE(weight, ' grams', '' ) AS NUMERIC)
        ELSE CAST(weight AS NUMERIC)
      END
    ) AS rn,
    COUNT(*) OVER() AS cnt
  FROM products
  WHERE weight IS NOT NULL
),
median_weight AS (
  SELECT ROUND(AVG(CAST(weight AS NUMERIC)), 2) AS median_value
  FROM ranked_weights
  WHERE rn IN ((cnt + 1) / 2, (cnt + 2) / 2)
),
median_price AS (
  SELECT ROUND(AVG(price::NUMERIC), 2) AS median_value
  FROM (
    SELECT price, ROW_NUMBER() OVER (ORDER BY price::NUMERIC) AS rn, COUNT(*) OVER() AS cnt
    FROM products WHERE price IS NOT NULL
  ) sub
  WHERE rn IN ((cnt + 1) / 2, (cnt + 2) / 2)
)
SELECT
  product_id,
  COALESCE(NULLIF(product_type, ''), 'Unknown') AS product_type,
  COALESCE(NULLIF(brand, '-'), 'Unknown') AS brand,
  ROUND(
    COALESCE(
      CASE
        WHEN weight ~ E'^\\d+\\.?\\d*\\s*grams$' THEN CAST(REPLACE(weight, ' grams', '' ) AS NUMERIC)
        ELSE CAST(weight AS NUMERIC)
      END,
      (SELECT median_value FROM median_weight)
    ), 2
  ) AS weight,
  ROUND(
    COALESCE(price::NUMERIC, (SELECT median_value FROM median_price)), 2
  ) AS price,
  COALESCE(average_units_sold::INTEGER, 0) AS average_units_sold,
  COALESCE(year_added::INTEGER, 2022) AS year_added,
  COALESCE(NULLIF(UPPER(stock_location), ''), 'Unknown') AS stock_location
FROM products;

```

...	↑↓	p...	...	↑↓	prod...	...	↑↓	brand	...	↑↓	...	↑↓	...	↑↓	average_units_...	...	↑↓	y...	...	↑↓	stock_lo...	...	↑↓
0				1	Bakery			TopBrand			602.61	11					15			2022			C
1				2	Produce			SilverLake			478.26	8.08					22			2022			C
2				3	Produce			TastyTreat			532.38	6.16					21			2018			B
3				4	Bakery			StandardYums			453.43	7.26					21			2021			D
4				5	Produce			GoldTree			588.63	7.88					21			2020			A
5				6	Meat			TopBrand			612.06	16.2					24			2017			A
6				7	Produce			GoldTree			320.49	8.01					21			2019			B
7				8	Meat			SilverLake			535.19	15.77					28			2021			A
8				9	Meat			StandardYums			375.07	11.57					30			2020			A
9				10	Meat			TastyTreat			506.34	13.94					27			2018			C
10				11	Dairy			StandardYums			345.07	9.26					26			2020			B
11				12	Bakery			StandardYums			345.58	6.87					21			2022			D
12				13	Snacks			SmoothTaste			512.54	8.65					19			2016			A
13				14	Meat			StandardYums			395.76	11.92					30			2019			A
14				15	Produce			SilverLake			324.92	7.94					23			2021			D
15				16	Dairy			SmoothTaste			446.76	10.79					23			2017			D

Rows: 1,700

[Expand](#)

Task 3

To find out how the range varies for each product type, your manager has asked you to determine the minimum and maximum values for each product type.

Write a query to return the `product_type`, `min_price` and `max_price` columns.

Unknown integration DataFrame as m

-- Write your query for task 3 in this cell

```
SELECT
    product_type,
    MIN(price::NUMERIC) AS min_price,
    MAX(price::NUMERIC) AS max_price
FROM products
GROUP BY product_type
ORDER BY product_type;
```

...	↑↓	prod...	...	↑↓	pr	...	↑↓	pr	...	↑↓
	0	Bakery			6.26			11.88		
	1	Dairy			8.33			13.97		
	2	Meat			11.48			16.98		
	3	Produce			3.46			8.78		
	4	Snacks			5.2			10.72		

Rows: 5

[Expand](#)

Task 4

The team want to look in more detail at meat and dairy products where the average units sold was greater than ten.

Write a query to return the `product_id`, `price` and `average_units_sold` of the rows of interest to the team.

Unknown integration DataFrame as

```
-- Write your query for task 4 in this cell
SELECT
  product_id,
  price::NUMERIC AS price,
  average_units_sold
FROM products
WHERE product_type IN ('Meat', 'Dairy')
  AND average_units_sold > 10
ORDER BY product_id;
```



...	↑↓	p...	...	↑↓	...	↑↓	average_units_...	...	↑↓
	0			6			16.2		24
	1			8			15.77		28
	2			9			11.57		30
	3			10			13.94		27
	4			11			9.26		26
	5			14			11.92		30
	6			16			10.79		23
	7			19			13.62		26
	8			20			13.03		22
	9			23			13.07		22
	10			24			10.98		23
	11			25			12.81		24
	12			28			13.01		20
	13			31			13.11		20
	14			41			8.63		27
	15			42			12.56		24

Rows: 698

Expand

FORMATTING AND NAMING CHECK

Use the code block below to check that your outputs are correctly named and formatted before you submit your project.

This code checks whether you have met our automarking requirements: that the specified DataFrames exist and contain the required columns. It then prints a table showing  for each column that exists and  for any that are missing, or if the DataFrame itself isn't available.

If a DataFrame or a column in a DataFrame doesn't exist, carefully check your code again.

IMPORTANT: even if your code passes the check below, this does not mean that your entire submission is correct. This is a check for naming and formatting only.

```

import pandas as pd

def check_columns(output_df, output_df_name, required_columns):
    results = []
    for col in required_columns:
        exists = col in output_df.columns
        results.append({'Dataset': output_df_name, 'Column': col, 'Exists': '✅' if exists else '❌'})
    return results

def safe_check(output_df_name, required_columns):
    results = []
    if output_df_name in globals():
        obj = globals()[output_df_name]
        if isinstance(obj, pd.DataFrame):
            results.extend(check_columns(obj, output_df_name, required_columns))
        elif isinstance(obj, str) and ("SELECT" in obj.upper() or "FROM" in obj.upper()):
            results.append({'Dataset': output_df_name, 'Column': '-', 'Exists': '💡 SQL query string'})
        else:
            results.append({'Dataset': output_df_name, 'Column': '-', 'Exists': '❌ Not a DataFrame or query'})
    else:
        results.append({'Dataset': output_df_name, 'Column': '-', 'Exists': '❌ Variable not defined'})
    return results

requirements = {
    'missing_year': ['missing_year'],
    'clean_data': ['product_id', 'product_type', 'brand', 'weight', 'price', 'average_units_sold', 'year_added',
    'stock_location'],
    'min_max_product': ['product_type', 'min_price', 'max_price'],
    'average_price_product': ['product_id', 'price', 'average_units_sold']
}

all_results = []
for output_df_name, cols in requirements.items():
    all_results += safe_check(output_df_name, cols)

check_results_df = pd.DataFrame(all_results)

print(check_results_df)

```

	Dataset	Column	Exists
0	missing_year	missing_year	✅
1	clean_data	product_id	✅
2	clean_data	product_type	✅
3	clean_data	brand	✅
4	clean_data	weight	✅
5	clean_data	price	✅
6	clean_data	average_units_sold	✅
7	clean_data	year_added	✅
8	clean_data	stock_location	✅
9	min_max_product	product_type	✅
10	min_max_product	min_price	✅
11	min_max_product	max_price	✅
12	average_price_product	product_id	✅
13	average_price_product	price	✅
14	average_price_product	average_units_sold	✅